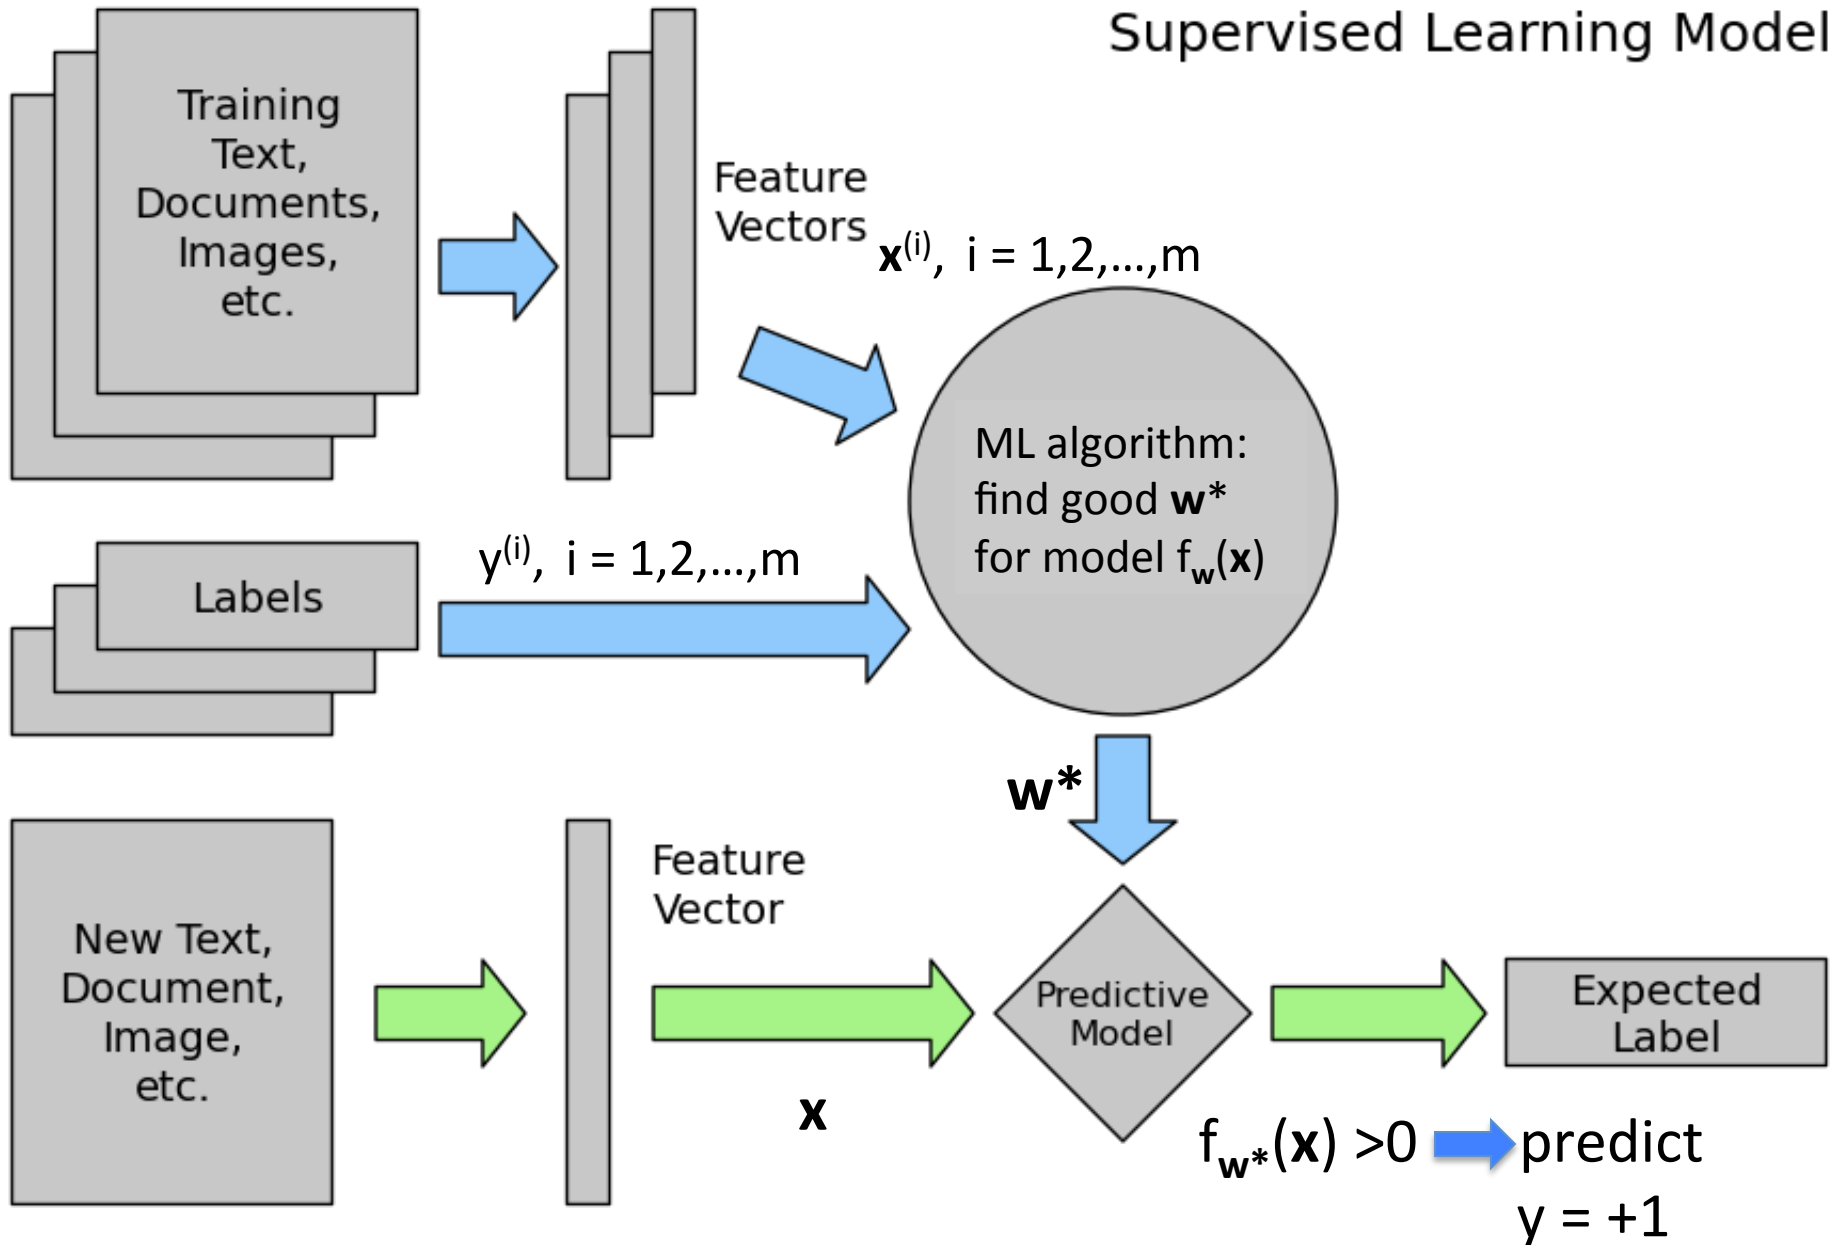


Supervised Machine Learning

An Overview of what
we've seen so far

Supervised Learning Model



Models so far

- Perceptron:
 - Linear model $h_{\mathbf{w}}(\mathbf{x})$
 - Requires data to be separable.
 - Finds good, not optimal \mathbf{w}^* using a simple algorithm.
- Hard-margin SVM:
 - Linear model $h_{\mathbf{w}}(\mathbf{x})$
 - Requires data to be separable.
 - Finds “optimal” \mathbf{w}^* by solving a constrained optimization problem to maximize margin
- Soft-margin SVM:
 - Linear model $h_{\mathbf{w}}(\mathbf{x})$
 - General data
 - Finds optimal \mathbf{w}^* by solving a constrained optimization problem
- Soft-margin SVM with kernels.
 - Non-linear model $h_{\mathbf{w}}(\mathbf{x})$
 - General data
 - Finds optimal \mathbf{w}^* by solving dual optimization problem. Uses SMO (Sequential Minimal Optimization)

Models so far (Part 2)

- Logistic Regression with MSE loss function:
 - Non-linear model $h_{\mathbf{w}}(\mathbf{x})$
 - But linear decision boundary
 - General data
 - Non-convex
 - Finds local optimal \mathbf{w}^* with gradient descent
- Logistic Regression with MLE loss function:
 - Convex
 - Finds optimal \mathbf{w}^* with gradient descent
- Logistic Regression with MLE loss function:
 - Convex
 - Finds optimal \mathbf{w}^* with gradient descent


Logistic Regression: MLE

- Training set: $\{(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}), \dots, (\mathbf{x}^{(m)}, y^{(m)})\}$
- How about choosing w to minimize MSE (as usual)?

$$J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\mathbf{w}}(\mathbf{x}^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\mathbf{w}}(\mathbf{x}), y) = \frac{1}{2} (h_{\mathbf{w}}(\mathbf{x}) - y)^2$$

$$h_{\mathbf{w}}(x) = \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}}}$$



Gradient Descent

$$J(\mathbf{w}) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\mathbf{w}}(\mathbf{x}^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\mathbf{w}}(\mathbf{x}^{(i)})) \right]$$

Want $\min_{\mathbf{w}} J(\mathbf{w})$:

Repeat {

$$w_j := w_j - \alpha \frac{\partial}{\partial w_j} J(\mathbf{w})$$

(simultaneously update all \mathbf{w}_j)

}

$$\frac{\partial}{\partial w_j} J(\mathbf{w}) = \frac{1}{m} \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$w_j := w_j - \alpha \sum_{i=1}^m (h_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$

Algorithm looks identical to linear regression! But it is not!

Gradient Descent Algorithm

Initialize \mathbf{w}

Repeat:

For $j = 0, 1, \dots, n$

$$\tilde{w}_j = w_j - \alpha \sum_{i=1}^m [f_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}] x_j^{(i)}$$

$\mathbf{w} \leftarrow \tilde{\mathbf{w}}$

Stochastic Gradient Descent (SGD)

Initialize \mathbf{w}

Repeat:

For $i = 1, 2, \dots, m$

For $j = 0, 1, \dots, n$

$$\tilde{w}_j = w_j - \alpha [f_{\mathbf{w}}(\mathbf{x}^{(i)}) - y^{(i)}] x_j^{(i)}$$

$$\mathbf{w} \leftarrow \tilde{\mathbf{w}}$$

Each update of \mathbf{w} uses a single data point $\mathbf{x}^{(i)}$!

General Gradient Descent

$$\mathbf{w} = \mathbf{w} - \alpha \nabla_{\mathbf{w}} J(\mathbf{w})$$

$$J(\mathbf{w}) = \sum_i \text{cost}(h_{\mathbf{w}}(\mathbf{x}^{(i)}), y^{(i)})$$

What is $\text{cost}(h_{\mathbf{w}}(\mathbf{x}), y)$ for MSE? For MLE?

What is a general expression for the partial derivative?

Models so far (3)

- Neural networks with MLE loss function:
 - Looking at binary classification for now
 - Use same MLE cost function as logistic regression
 - Finds \mathbf{w}^* with gradient descent
 - To evaluate $h_{\mathbf{w}}(\mathbf{x})$ use forward propagation
 - To evaluate partial derivatives, use back propagation
 - They are different from logistic regression. Why?
 - Let's look into the partial derivatives.